

神經網路與損失函數

黃志勝 (Tommy Huang)

義隆電子 人工智慧研發部

國立陽明交通大學 AI學院 合聘助理教授

國立台北科技大學 電資學院合聘助理教授



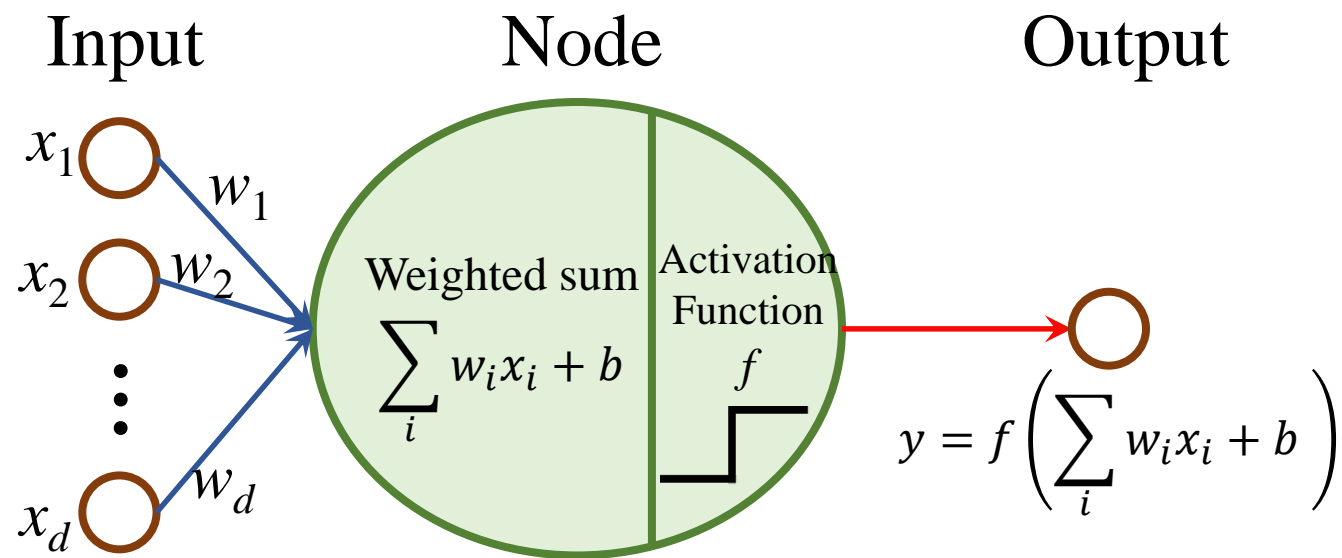


Outline

- 1. 類神經網路(Neural Network, NN)
- 2. 感知機(Perception)
- 3. Multi-layer perception (MLP)
- 4. How NN work?
- 5. Loss function



類神經網路

基本上神經網路是基於感知機(Perceptron)神經網路開始，主要是希望用數學模型去模擬神經細胞的運作模式。



權重(w_i): Dendrite




Input(x_i) and output (y) node: Synapse

Node: Cell body

Output: Axon






NN for classification problem




Perception

Perception can learn the nonlinear representation by activation function.



Perception



$$y = f(w_{10} + w_{11}x_1 + w_{12}x_2 + \cdots + w_{1d}x_d) = f(\mathbf{W}^T \mathbf{x})$$

$$\text{Classification: } f = \begin{cases} 1 & \mathbf{W}^T \mathbf{x} \geq 0 \\ 0 & \text{O.W.} \end{cases}$$


$$\text{Regression: } f(\mathbf{W}^T \mathbf{x})$$

$$\mathbf{W} = \begin{bmatrix} w_0 \\ w_1 \\ \vdots \\ w_d \end{bmatrix}, \mathbf{x} = \begin{bmatrix} 1 \\ x_1 \\ \vdots \\ x_d \end{bmatrix}$$



神經元與回歸

神經元



Perception with linear output is the linear regression.

NN: backpropagation

Regression: OLSE (ordinary least squares estimator).

回歸

$$f(\mathbf{x}) = \beta_0 + \beta_1 x_1 + \cdots + \beta_d x_d$$


$$y = \sigma(f(\mathbf{x})) = \frac{1}{1 + e^{-f(\mathbf{x})}} = \frac{1}{1 + e^{-\beta^T \mathbf{x}}}$$




NN for XOR problem

- Exclusive OR (XOR Boolean function)
- It's impossible to find a single straight line to separate two classes.

Truth Table for the XOR problem			
x1	x2	XOR	Class
0	0	0	B
0	1	1	A
1	0	1	A
1	1	0	B



NN for XOR problem



OR


$$h_1(x) = x_1 + x_2 - 0.5 = 0$$

AND

$$h_2(x) = x_1 + x_2 - 1.5 = 0$$



NN for XOR problem



OR


$$h_1(x) = x_1 + x_2 - 0.5 = 0$$

AND

$$h_2(x) = x_1 + x_2 - 1.5 = 0$$



NN for XOR problem



OR


$$h_1(x) = x_1 + x_2 - 0.5 = 0$$

AND


$$h_2(x) = x_1 + x_2 - 1.5 = 0$$



NN for XOR problem



從原始特徵空間做特徵轉換(特徵萃取)




Truth Table for XOR problem

x1	x2	$h_1(x)$	$h_2(x)$	Class
0	0	0(-)	0(-)	B
0	1	1(+)	0(-)	A
1	0	1(+)	0(-)	A
1	1	1(+)	1(+)	B



Two Layer Perception

$$y_1 = f(h_1(x)) = f(x_1 + x_2 - 0.5)$$



$$y_2 = f(h_2(x)) = f(x_1 + x_2 - 1.5)$$


Step activation function

$$f(v) = \begin{cases} 1 & \text{if } v \geq 0 \\ 0 & \text{if } v < 0 \end{cases}$$




類神經網路

但神經訊息傳遞不會像只有一層Single layer perceptron運作，神經網路應該是多個細胞部段將訊息傳遞下去運作的模式，這就是Multilayer perception (MLP)，也就是一般認知的類神經網路。



MLP




在深度學習這樣的層跟層叫做fully connection



Activation

- 如果沒有activation function會有什麼影響



$$X_2 = W_1 X_1$$

$$X_3 = W_2 X_2$$

$$X_4 = W_3 X_3$$

$$Y = W_4 X_4$$

$$Y = W_4 W_3 W_2 W_1 X_1$$

$$\underline{Y = W X_1}$$


$$W = W_4 W_3 W_2 W_1$$

結果就只是一層的神經網路而已。



Activation

- 如果沒有activation function會有什麼影響



$$\begin{aligned} X_2 &= W_1 X_1 \\ X_3 &= W_2 X_2 \\ X_4 &= W_3 X_3 \\ Y &= W_4 X_4 \end{aligned}$$



$$\begin{aligned} X_2 &= f(W_1 X_1) \\ X_3 &= f(W_2 X_2) \\ X_4 &= f(W_3 X_3) \\ Y &= f(W_4 X_4) \end{aligned}$$

$$Y = W_4 W_3 W_2 W_1 X_1 \quad Y = f(W_4 f(W_3 f(W_2 f(W_1 X_1))))$$

$Y = W X_1$

$$W = W_4 W_3 W_2 W_1$$



Activation Function

- Sigmoid 、Tanh 。

ReLU(Rectified Linear Unit)系列：


- ReLU
- Leaky ReLU
- ReLU6
- PReLU
- RReLU
- ELU (Exponential Linear Unit)
- SELU (Scaled Exponential Linear Units)



Activation function


- $\text{Sigmoid}(x) = \frac{1}{1+e^{-x}}$

將值壓到0~1之間




- $\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$

將值壓到-1~1之間



Sigmoid近似法

$$y = \text{approxSigmoid}(x) = \begin{cases} 0.5 * \left(1.5 * \left(\frac{\frac{x}{2}}{1 + \frac{x}{2}} \right) + 1 \right) & \text{if } 0 \leq x < 3.4 \\ 0.5 * (0.9444 + 0.0459 * (\frac{x}{2} - 1.7) + 1) & \text{if } 3.4 \leq x \leq 5.8 \\ 0.9997 & x \geq 5.8 \end{cases}$$



Black: sigmoid
Red: approximate sigmoid

```
def fast_sigmoid2(x):
    x = 0.5 * x
    flag_neg=0
    if x<0:
        x=-x
        flag_neg=1

    if x < 1.7:
        z = (1.5 * x / (1 + x))
    elif x < 2.9:
        z = (0.9444 + 0.0459 * (x - 1.7))
    else:
        z = 0.9997
    if flag_neg==1:
        z=-z
    return 0.5 * (z + 1.)
```



Activation function: ReLU系列


- $\text{ReLU}(x) = \max(0, x) = \begin{cases} x & \text{if } x > 0 \\ 0 & \text{O.W.} \end{cases}$
- $\text{ReLU6}(x) = \min(\max(0, x), 6) = \begin{cases} 6 & \text{if } x \geq 6 \\ x & \text{if } 0 < x < 6 \\ 0 & \text{O.W.} \end{cases}$
- $\text{LeakyReLU}(x) = \max(0, x) + a * \min(0, x) = \begin{cases} x & \text{if } x > 0 \\ ax & \text{O.W.} \end{cases}$ (default $a = 0.1$)
- $\text{PReLU}(x) = \max(0, x) + a * \min(0, x) = \begin{cases} x & \text{if } x \geq 0 \\ ax & \text{O.W.} \end{cases}$ (a 是訓練得到, init 設定: 0.25)
- $\text{RReLU}(x) = \max(0, x) + a * \min(0, x) = \begin{cases} x & \text{if } x > 0 \\ ax & \text{O.W.} \end{cases}$ (a 是隨機 $U(\text{lower} = \frac{1}{8}, \text{upper} = \frac{1}{3})$ 選取)
- $\text{ELU}(x) = \max(0, x) + \min(0, \alpha * (e^x - 1)) = \begin{cases} x & \text{if } x > 0 \\ \alpha * (e^x - 1) & \text{O.W.} \end{cases}$
- $\text{SELU}(x) = \text{scale}(\max(0, x) + \min(0, \alpha * (e^x - 1))) = \text{scale} \begin{cases} x & \text{if } x > 0 \\ \alpha * (e^x - 1) & \text{O.W.} \end{cases}$
 $\alpha=1.6732632423543772848170429916717$
 $\text{scale}=1.0507009873554804934193349852946$
 實驗驗證, 小數取兩位數即可, $\alpha=1.67$, $\text{scale}=1.05$



SELU係數精度實驗

- Model: LeNet 加強版
- Database: Cifar-10
- 看分類正確率隨著learning epoch變化的影響

SELU係數小數位數全取





$$\text{SELU}(x) = \text{scale}(\max(0, x) + \min(0, \alpha * (e^x - 1)))$$

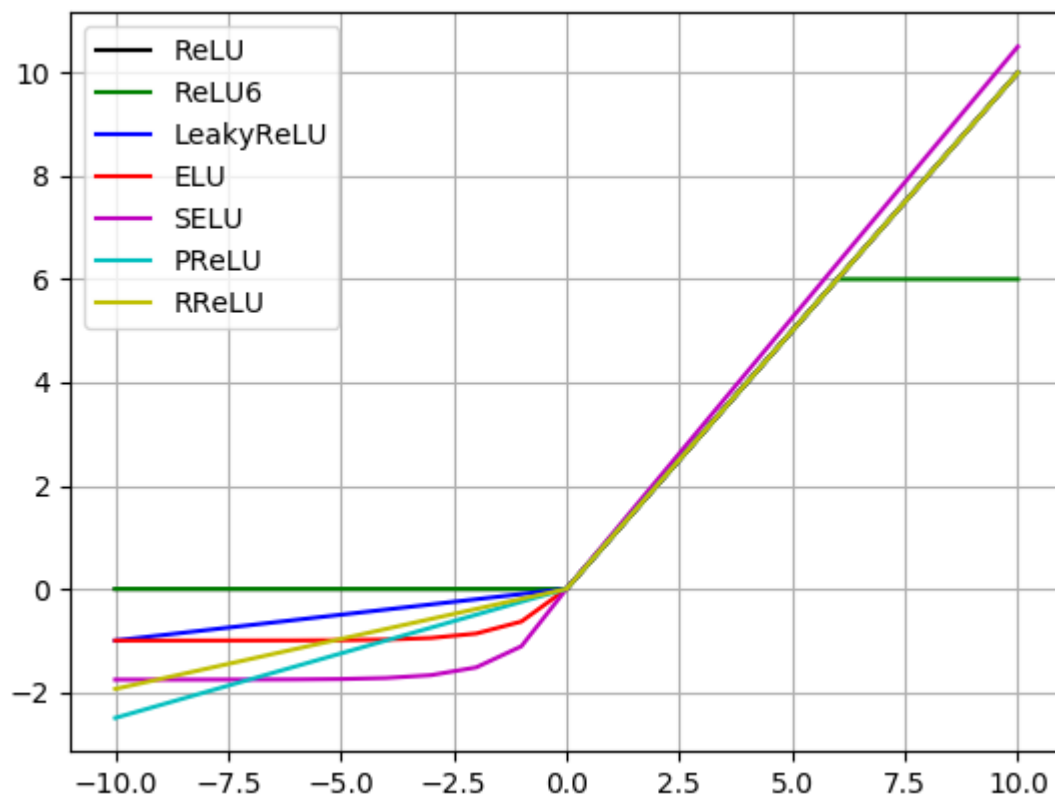
$$= \text{scale} \begin{cases} x & \text{if } x > 0 \\ \alpha * (e^x - 1) & \text{O.W.} \end{cases}$$

$\alpha=1.6732632423543772848170429916717$


$\text{scale}=1.0507009873554804934193349852946$



Activation function: ReLU系列





Polyhedral Regions




Polyhedral Regions

activation function = step function



The first layer of neurons divides the input d-dimensional space into **polyhedral**, which are formed by hyperplane intersections.




Polyhedral Regions




All vectors located within one of these polyhedral regions are mapped onto a specific vertex of the unit hypercube.




Dimension Reduction (Feature Projection)




Principle Component Analysis



PCA




NN and Dimension Reduction



$$\mathbf{x}' = \begin{bmatrix} x'_1 \\ x'_2 \end{bmatrix} = \begin{bmatrix} w_{11} & w_{12} \\ w_{21} & w_{22} \end{bmatrix}^T \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \mathbf{W}^T \mathbf{x}$$

拓展
為d維




$$\mathbf{x}' = \begin{bmatrix} x'_1 \\ \vdots \\ x'_d \end{bmatrix} = \mathbf{W}^T \mathbf{x} = \begin{bmatrix} w_{11} & \cdots & w_{1d} \\ \vdots & \ddots & \vdots \\ w_{d1} & \cdots & w_{dd} \end{bmatrix}^T \begin{bmatrix} x_1 \\ \vdots \\ x_d \end{bmatrix}$$



NN and Dimension Reduction

- Principle Component Analysis

$$\mathbf{x}' = \begin{bmatrix} x_1' \\ \vdots \\ x_d' \end{bmatrix} = W^T \mathbf{x} = \begin{bmatrix} w_{11} & \dots & w_{1d} \\ \vdots & \ddots & \vdots \\ w_{d1} & \dots & w_{dd} \end{bmatrix} \begin{bmatrix} x_1 \\ \vdots \\ x_d \end{bmatrix}$$




PCA是NN的一個special case

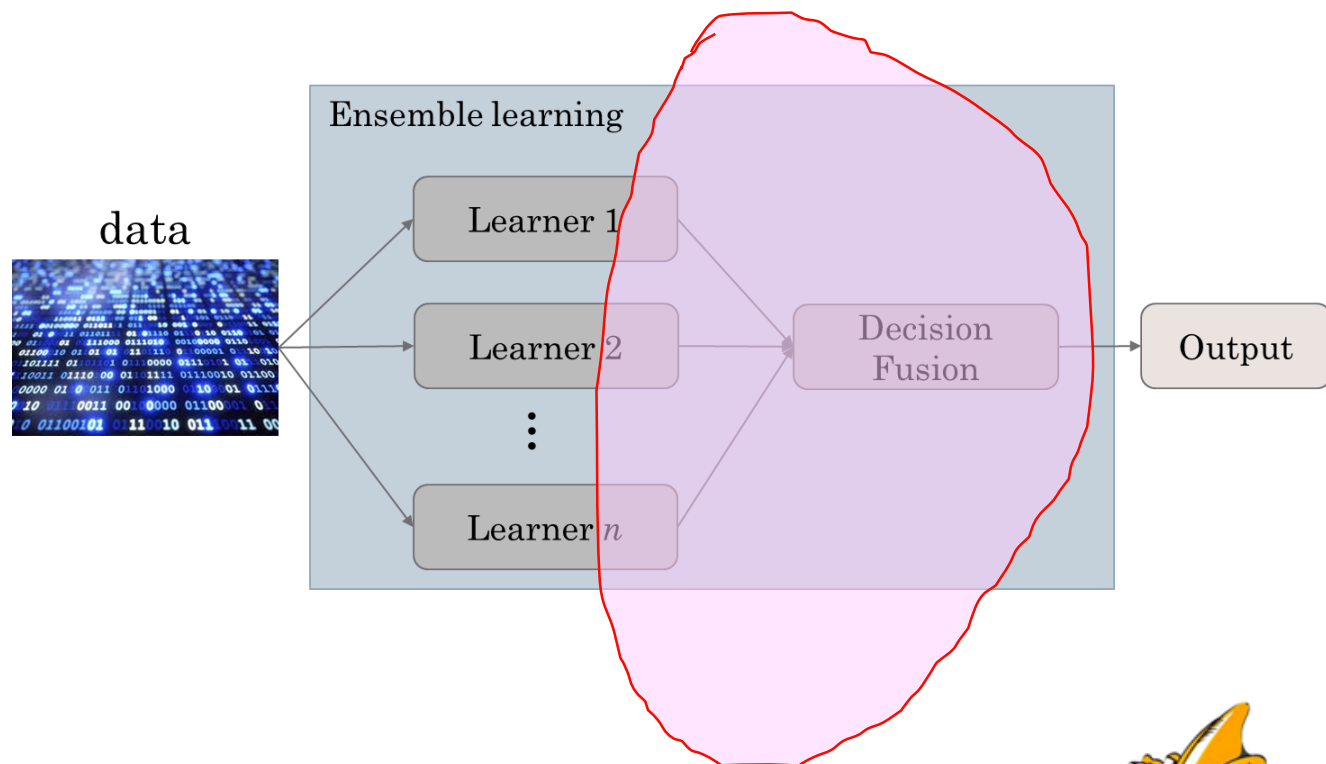


NN and Ensemble Learning

NN



Ensemble Learning



Introduction (loss function)

In learning algorithm, there is an assumption, which may accompany with an **object function**.

Regression: minimizing the mean square errors (MSE)

K-mean: minimizing the mean of square error between data and centers.

PCA: maximizing the variance of project data.

SVM: maximizing the margin.

Sometimes the same model but different object function can lead different results. (Linear regression and ridge regression)

Regression

$$\min_{\beta} \{ \text{MSE}(\hat{y}, y) \}$$

Ridge regression

$$\min_{\beta} \{ \text{MSE}(\hat{y}, y) + \lambda L_2 \text{norm}(\beta) \}$$

$$L_2 \text{norm}(\beta) = \sum_i \beta_i^2$$





Neural Network

(Number of output node)

Regression case: Number of output node of NN architecture is 1.

Classification case: Number of output node of NN architecture is L (Number of class in define problem). -> **memberships (call logits in DL) for each class (such as posterior probability in ML).**

Q: Why number of output node in this toy case is 1?



ANS: It's a binary classification problem.


A simple threshold can be defined as decision boundary, see the Logistic regression.




Neural Network

(Number of output node)


羅吉斯迴歸



$$\sigma(f(x)) = \frac{1}{1 + e^{-f(x)}} = \frac{e^{f(x)}}{1 + e^{f(x)}}$$



Neural Network



Logistic Regression	NN	Activation
$f(x)$	$h(x)$	None
$\sigma(f(x))$	$f(h(x))$	Sigmoid



Introduction

1. Regression

MSE, MAE, Huber Loss

2. Classification

Cross entropy, Focal loss

3. Triple loss



Residual

- Residual: predicted value v.s. target value.

Regression:

$$y - \hat{y}$$

Classification (error):

$$sign(\hat{y}, y) = \begin{cases} 1 & \hat{y} = y \\ 0 & \hat{y} \neq y \end{cases}$$

$$error\ rate = \frac{1}{n} \sum_{i=1}^n sign(\hat{y}_i, y_i)$$



MSE & MAE

Mean Square Error (MSE)

Mean Absolute Error (MAE)

Why square or absolute?

Target value: $y_1 = 0, y_2 = 1$

Predicted value: $\hat{y}_1 = 100, \hat{y}_2 = 99$

$$\text{Residual 1} = y_1 - \hat{y}_1 = 0 - 100 = -100$$

$$\text{Residual 2} = y_2 - \hat{y}_2 = 1 - (-99) = 100$$

$$\text{Residual 1} + \text{Residual 2} = -100 + 100 = 0$$



MSE & MAE

Mean Square Error (MSE)

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$


Mean Absolute Error (MAE)

$$MAE = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$



MSE & MAE

Trend of residual (target value =0)



Comparison (MSE&MAE)

With a fair baseline, RMSE (root MSE)

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

$$RMSE = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$



Comparison (MSE&MAE)

Change ID5 with a outlier

ID	residual	residual	residual ²
1	-10	10	100
2	-5	5	25
3	0	0	0
4	5	5	25
5	10	10	100
MAE=6, RMSE=7.07			

ID	residual	residual	residual ²
1	-10	10	100
2	-5	5	25
3	0	0	0
4	5	5	25
5	100	100	10000
MAE=24, RMSE=45.06			

Problem of MSE: more outlier sensitivity.



Comparison (MSE&MAE)

- Problem of MAE: same gradient value.
- When loss is small, it's difficult to reach the optimal target.

$$\begin{aligned} f_1(x) &= x^2, f_1'(x) = 2x \\ f_2(x) &= |x|, f_2'(x) = \frac{x}{|x|} \end{aligned}$$


Gradient update:

$$x^{t+1} \rightarrow x^t - r f'(x)$$

Suppose: $x^0 = 2, r = 0.3$



Comparison (MSE&MAE)



$f(x)=x^2$				$f(x)= x $			
t	x^t	$f'(x)$	x^{t+1}	t	x^t	$f'(x)$	x^{t+1}
1	2	4	0.8	1	2	1	1.7
2	0.8	1.6	0.32	2	1.7	1	1.4
3	0.32	0.64	0.128	3	1.4	1	1.1
4	0.128	0.256	0.0512	4	1.1	1	0.8
5	0.0512	0.1024	0.02048	5	0.8	1	0.5
6	0.02048	0.04096	0.008192	6	0.5	1	0.2
7	0.008192	0.016384	0.003277	7	0.2	1	-0.1
8	0.003277	0.006554	0.001311	8	-0.1	-1	0.2
9	0.001311	0.002621	0.000524	9	0.2	1	-0.1
10	0.000524	0.001049	0.00021	10	-0.1	-1	0.20
11	0.00021	0.000419	0.000084	11	0.20	1	-0.1
12	0.000084	0.000168	0.000034	12	-0.10	-1	0.20
13	0.000034	0.000067	0.000013	13	0.20	1	-0.10
14	0.000013	0.000027	0.000005	14	-0.10	-1	0.20
15	0.000005	0.000011	0.000002	15	0.20	1	-0.10
16	0.000002	0.000004	0.000001	16	-0.10	-1	0.20
17	0.000001	0.000002	0.000000	17	0.20	1	-0.10
18	0.000000	0.000001	0.000000	18	-0.10	-1	0.20
19	0.000000	0.000000	0.000000	19	0.20	1	-0.10
20	0.000000	0.000000	0.000000	20	-0.10	-1	0.20




Huber Loss

Huber loss:

$$Loss(y, \hat{y})$$

$$= \begin{cases} \frac{1}{2} (y - \hat{y})^2, & |y - \hat{y}| \leq \delta \\ \delta (|y - \hat{y}| - \frac{1}{2} \delta), & o.w. \end{cases}$$

δ : parameter of Huber loss.



MAE, MSE & Huber Loss

ID	residual	residual	residual ²	Huber ($\delta=1$)	Huber ($\delta=10$)
1	-10	10	100	9.5	50
2	-5	5	25	4.5	12.5
3	0	0	0	0	0
4	5	5	25	4.5	12.5
5	10	10	100	9.5	50
MAE=6, RMSE=7.07 MeanHuber($\delta=1$)=5.6, MeanHuber($\delta=10$)=25					

ID	residual	residual	residual ²	Huber ($\delta=1$)	Huber ($\delta=10$)
1	-10	10	100	9.5	50
2	-5	5	25	4.5	12.5
3	0	0	0	0	0
4	5	5	25	4.5	12.5
5	100	100	10000	99.5	950
MAE=24, RMSE=45.06 MeanHuber($\delta=1$)=23.6, MeanHuber($\delta=10$)=205					



Classification

Classification:

$$\text{sign}(\hat{y}, y) = \begin{cases} 1 & y = \hat{y} \\ 0 & y \neq \hat{y} \end{cases}$$
$$\text{error rate} = 1 - \frac{1}{n} \sum_{i=1}^n \text{sign}(\hat{y}_i, y_i)$$

We hope less error rate more better in classification.

Can we use the classification error rate/accuracy as loss function?



Classification

		Model 1 (輸出)				Model 2 (輸出)			
		機率輸出			判斷	機率輸出			判斷
	Target (Label)	男生	女生	其他		男生	女生	其他	
data 1	男生	0.4	0.3	0.3	男生 (正確)	0.7	0.1	0.2	男生 (正確)
data 2	女生	0.3	0.4	0.3	女生 (正確)	0.1	0.8	0.1	女生 (正確)
data 3	男生	0.5	0.2	0.3	男生 (正確)	0.9	0.1	0	男生 (正確)
data 4	其他	0.8	0.1	0.1	男生 (錯誤)	0.4	0.3	0.3	男生 (錯誤)
		模型1錯誤率: $1/4=0.25$				模型2錯誤率: $1/4=0.25$			

Can we observe any difference between model 1 & 2 from error rate?


NO...

BUT we can observe that model 2 has better probability outputs than model 1.
Error rate cannot as learning object for learning updating, it's just a metric for evaluating model performance.



Classification

- How do we make decision for a new sample in classification model?
- **ANS: posterior probability.**



Cross-entropy

- Cross-entropy is usually used in classification loss.
- **Entropy**: the average of information which is produced by a stochastic source of data.
- **Information gain**: (suppose X is a random variable)

$$I(x) = -\log_2(p(x))$$



Information gain

A is stupid, and his grades usually are around 50 marks.

B is smart, and his grades usually are almost 100 marks.

Probability to pass the exam for A: $p(x_A) = 0.4$

$$I(x_A) = -\log_2(p(x_A)) = 1.322$$

Probability to pass the exam for B: $p(x_B) = 0.99$

$$I(x_B) = -\log_2(p(x_B)) = 0.014$$



Entropy

Entropy: the average of information which is produced by a stochastic source of data.

In information theory,

Entropy = Shannon entropy

$$H(X) = \sum_i -p_i \log_2(p_i)$$

Generally, entropy refers to uncertainty for the random variable X .



Entropy

$$p(x_A = \text{pass}) = 0.4, \quad p(x_A = \text{fail}) = 0.6$$

$$p(x_B = \text{pass}) = 0.99, \quad p(x_B = \text{fail}) = 0.01$$

$$H(X) = \sum_i -p_i \log_2(p_i)$$

$$H(X_A) = -0.4 \log(0.4) - 0.6 \log(0.6) = 0.971$$

$$H(X_B) = -0.99 \log(0.99) - 0.01 \log(0.01) = 0.081$$

Same conclusion for information gain.


$$I(x_A) = -\log_2(p(x_A)) = 1.322$$

$$I(x_B) = -\log_2(p(x_B)) = 0.014$$





Entropy

When $p=0.5$ has the largest entropy.



Entropy



$p(X_A) = \begin{cases} 0.1 & x = 1 \\ 0.15 & x = 2 \\ 0.5 & x = 3 \\ 0.15 & x = 4 \\ 0.1 & x = 5 \end{cases}$	$p(X_A) = \begin{cases} 0.01 & x = 1 \\ 0.09 & x = 2 \\ 0.8 & x = 3 \\ 0.09 & x = 4 \\ 0.01 & x = 5 \end{cases}$	$p(X_B) = \begin{cases} 0.2 & x = 1 \\ 0.2 & x = 2 \\ 0.2 & x = 3 \\ 0.2 & x = 4 \\ 0.2 & x = 5 \end{cases}$
$H(X_A) = 1.985$	$H(X_A) = 1.016$	$H(X_B) = 2.322$



Cross-entropy

Formula of cross- entropy:

$$CE = \sum_{i=1}^n \sum_{c=1}^C -y_{c,i} \log_2(p_{c,i})$$

C : number of class (male, female, other)

n : number of data


$y_{c,i}$: binary indicator (0 or 1) from one hot encode (i -th data assigns to c -class)

$p_{c,i}$: probability of i -th data assigns to c -class




One hot encode (Dummy variable)

Data 1	Male
Data 2	Female
Data 3	Male
Data 4	Other



Male
Female
Other



	Male	Female	Other
Data 1	1	0	0
Data 2	0	1	0
Data 3	1	0	0
Data 4	0	0	1



Cross-entropy

$$CE = \sum_{i=1}^n \sum_{c=1}^C -y_{c,i} \log_2(p_{c,i})$$

		Model 1 (輸出)					
		機率輸出			實際One-hot encode		
	Target (Label)	男生	女生	其他	男生	女生	其他
data 1	男生	0.4	0.3	0.3	1	0	0
data 2	女生	0.3	0.4	0.3	0	1	0
data 3	男生	0.5	0.2	0.3	1	0	0
data 4	其他	0.8	0.1	0.1	0	0	1
		模型1錯誤率: 1/4=0.25 Cross-entropy=6.966					

Data 1:

$$\sum_{c=1}^C -y_{c,1} \log_2(p_{c,1})$$

$$= -1 * \log(0.4) - 0 * \log(0.3) - 0 * \log(0.3) = 1.322$$

Data 4:

$$\sum_{c=1}^C -y_{c,4} \log_2(p_{c,4})$$

$$= -0 * \log(0.8) - 0 * \log(0.1) - 1 * \log(0.1) = 3.3219$$

SO less probability data has larger loss function (entropy value)→learning target.




Cross-entropy for evaluating the model performance

		Model 1 (輸出)					
		機率輸出			實際One-hot encode		
	Target (Label)	男生	女生	其他	男生	女生	其他
data 1	男生	0.4	0.3	0.3	1	0	0
data 2	女生	0.3	0.4	0.3	0	1	0
data 3	男生	0.5	0.2	0.3	1	0	0
data 4	其他	0.8	0.1	0.1	0	0	1
		模型1錯誤率: $1/4=0.25$ Cross-entropy=6.966					

		Model 2 (輸出)					
		機率輸出			實際One-hot encode		
	Target (Label)	男生	女生	其他	男生	女生	其他
data 1	男生	0.7	0.1	0.2	1	0	0
data 2	女生	0.1	0.8	0.1	0	1	0
data 3	男生	0.9	0.1	0	1	0	0
data 4	其他	0.4	0.3	0.3	0	0	1
		模型1錯誤率: $1/4=0.25$ Cross-entropy= 2.310					

Cross-entropy for loss function



$$CE\ loss = \sum_{i=1}^n \sum_{c=1}^C -y_{c,i} \log_2(p_{c,i})$$




Focal loss (1/3)

Cross-entropy (CE) for $y \in \{\pm 1\}$

$$CE(p, y) = \begin{cases} -\log(p), & \text{if } y = 1 \\ -\log(1 - p), & \text{if } y = -1 \end{cases}$$

$$CE(p, y) = CE(p_t) = -\log(p_t), \quad p_t = \begin{cases} p, & \text{if } y = 1 \\ 1 - p, & \text{if } y = -1 \end{cases}$$



Focal loss (2/3)

α -balanced cross-entropy:

$$CE(p_t) = -\alpha \log(p_t)$$

BUT it's not effect for larger class unbalance problem.

Modulating factor:

$$(1 - p_t)^r$$

r : focusing parameter, $r \geq 0$.

Focal loss:


$$FL(p_t) = -(1 - p_t)^r \log(p_t)$$

α -balanced focal loss:

$$FL(p_t) = -\alpha(1 - p_t)^r \log(p_t)$$



Focal loss (3/3)



Why Focal loss? (Proposed in RetainNet)


在訓練階段的特徵圖對應的Anchor
有90%都是沒用的背景anchor，所以要怎麼抑制背景的anchor。

怎麼強化判錯的Anchor的loss?


Focal loss。



3*3的特徵圖



7*7的特徵圖




Triple Loss

Triple

Anchor: a randomly training data with label c

Positive: training data in label c

Negative: training data in other labels



Triple Loss

Anchor : x_i^a

Positive : x_i^p

Negative : x_i^n

Encoder network : $f(x)$

Anchor : $f(x_i^a)$


Positive : $f(x_i^p)$

Negative : $f(x_i^n)$

triple loss aims to

$\text{dist}(f(x_i^a), (x_i^p)) \downarrow$

$\text{dist}(f(x_i^a), (x_i^n)) \uparrow$



Triple Loss

$$\text{dist}(f(x_i^a), f(x_i^p)) + \alpha < \text{dist}(f(x_i^a), f(x_i^n))$$

$$\Rightarrow \|f(x_i^a) - f(x_i^p)\|_2^2 + \alpha < \|f(x_i^a) - f(x_i^n)\|_2^2$$

$$\arg \min \left\{ \sum_i \left(\|f(x_i^a) - f(x_i^p)\|_2^2 + \alpha - \|f(x_i^a) - f(x_i^n)\|_2^2 \right) \right\}$$




$$\arg \min \sum_i \left[\|f(x_i^a) - f(x_i^p)\|_2^2 + \alpha - \|f(x_i^a) - f(x_i^n)\|_2^2 \right]_+$$



Triple Loss

$$\left[d_p - d_n + \alpha \right]_+ = \begin{cases} d_p - d_n + \alpha & d_p - d_n + \alpha > 0 \\ 0 & d_p - d_n + \alpha < 0 \end{cases}$$





Conclusion

MSE, MAE, Huber loss, triple loss do the same thing.

Similarity measurement.

Cosine loss



Can MSE be a loss for classification?

